



ENV797 - TIME SERIES ANALYSIS FOR ENERGY AND ENVIRONMENT APPLICATIONS

M6 - Seasonal ARIMA Models

Prof. Luana Medeiros Marangon Lima, Ph.D.

Learning goals



- Discuss the seasonal ARIMA model – SARIMA
- Learn how to identify order of seasonal component
- Learn how to fit SARIMA models in R

Seasonal Models

SARIMA


Modeling with Seasonality

- So far we have two ways to deal with seasonality
 - Linear regression on seasonal dummy variables and work with residuals
 - Computing and subtracting the average of each “month” (season)
- In R: use `stats::decompose()` and `forecast::seasadj()`
- Drawbacks
 - Need to add seasonality back after fitting the model
 - Seasonal component is modeled as “constant” over time



Seasonal Models

- Today we will learn a new approach: work with seasonal models, i.e., models that can handle seasonal time series
- We still need to identify the seasonal component but we do not need to remove it from the series



- And we don't need to add it back after fitting the model

Seasonal ARIMA Models

- The Seasonal ARIMA models rely on seasonal lags and difference to fit the seasonal pattern
- The seasonal part of the ARIMA model have three additional parameters

P: number of seasonal autoregressive terms

D: number of seasonal differences

Q: number of moving average terms

- The complete model is called an

$$\text{ARIMA}(p, d, q) \times (P, D, Q)_s$$

Non-seasonal
part

Seasonal part

Seasonal ARIMA Models (cont'd)

- We have two terms that refer to differencing
- How does all this differencing work?

$$\text{If } d = 0, D = 1: \quad y_t = Y_t - Y_{t-s}$$

$$\text{If } d = 1, D = 1: \quad y_t = (Y_t - Y_{t-1}) - (Y_{t-s} - Y_{t-s-1})$$

- Note that s is the seasonal period, i.e., for monthly data $s = 12$
- D should never be more than 1
- $(d + D)$ should never be more than 2

Seasonal ARIMA Models (cont'd)

- What about the seasonal AR and MA terms?

If $P = 1$: y_{t-s} is added to the equation P always refer to autoregressive

If $Q = 1$: a_{t-s} is added to the equation Q is the moving average.

- $(P + Q)$ should not be more than one

- **After differencing** plot ACF and PACF

- ▣ Positive spikes in ACF at lag 12, 24, 36, ...

$P = 1$

- and single positive spike in PACF at lag 12

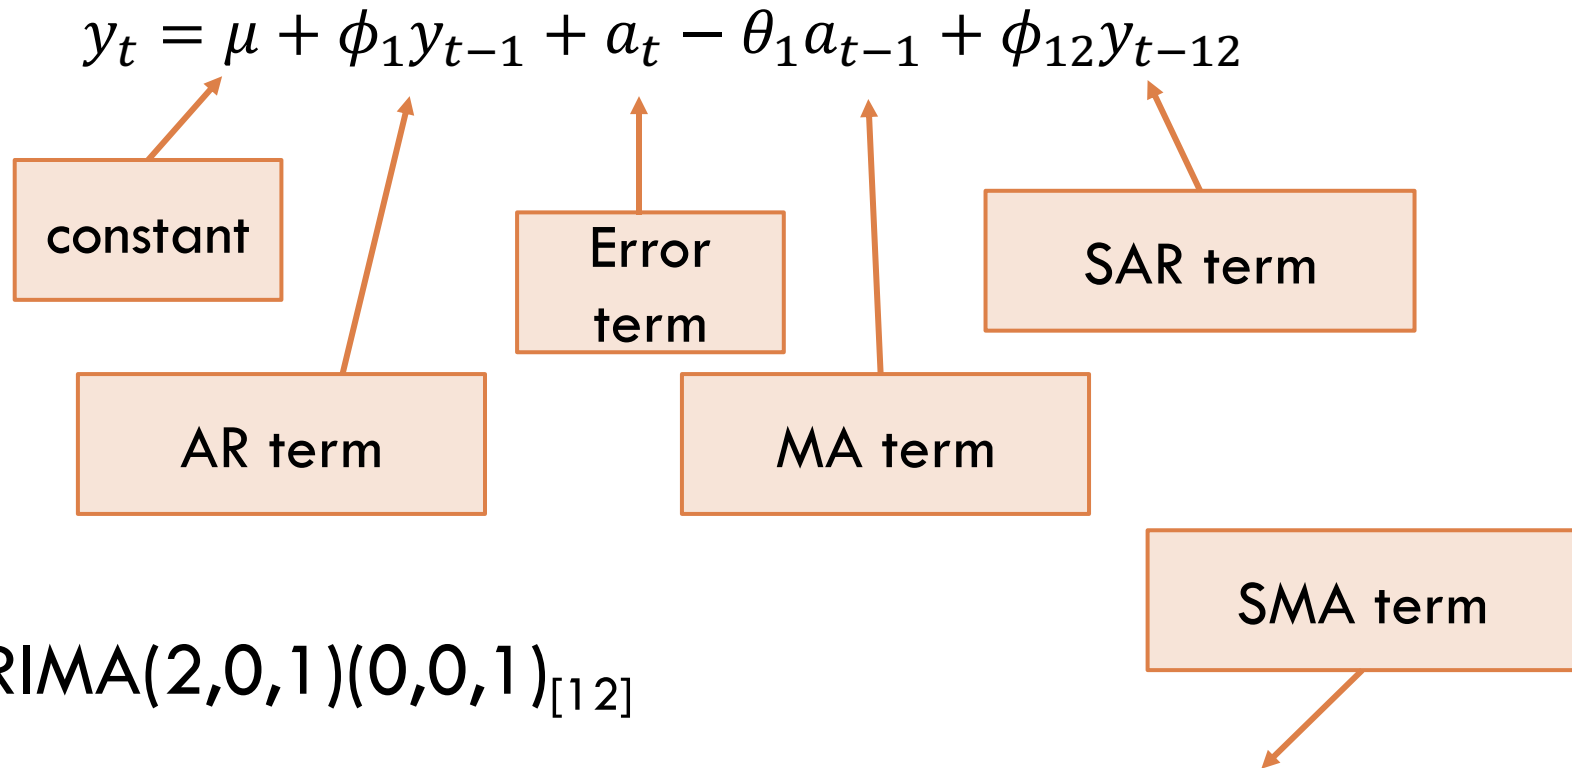
- ▣ Negative spike in ACF at lag 12

$Q = 1$

- and negative spikes in PACF at lag 12, 24, 36, ...

Examples: Seasonal ARIMA

- SARIMA(1,0,1)(1,0,0)_[1 2] If you difference once, you can have a drift, or no drift otherwise.



- SARIMA(2,0,1)(0,0,1)_[1 2]

$$y_t = \mu + \phi_1 y_{t-1} + \phi_2 y_{t-2} + a_t - \theta_1 a_{t-1} - \theta_{12} a_{t-12}$$

Periodic AR Model (PAR)

Fits a model for each season of the year for any order of choice.

- Also known as PAR(p) where p is the order of autoregressive model
- Similar to having one AR model for each season of the year

$$y_t = \phi_{1,s}y_{t-1} + \cdots + \phi_{p,s}y_{t-p} + a_t$$

- Note that $s = 1, 2, 3, \dots$ depending on the number of seasons
- The autoregressive parameter vary with the season for lag



ARIMA Models Summary

ARIMA Models Summary

- $AR(p), MA(q), ARMA(p,q)$ – stationary and non-seasonal process
- $ARIMA(p,d,q)$ – non-stationary and non-seasonal process
 - ▣ The Integrated part can handle the non-stationarity
 - ▣ The non-stationary could be either a deterministic or stochastic trend eliminated by differencing
- Seasonal ARIMA or SARIMA – non-stationary and seasonal process

No s for this

$$ARIMA(p, d, q) \times (P, D, Q)_s$$

Non-seasonal
part

Seasonal part

Fitting ARIMA Models

Cheat sheet

How to Specify Model Orders?

Start with non-seasonal part


- Step 1: Run stationary test
 - To identify stochastic trend use ADF
 - To identify determinist trend use Mann Kendall
- Step 2: If either trend is present $d = 1$
 - It may be the case that you need differencing more than once to remove trend
 - To find out if $d = 1$ is enough, run the tests again on the differenced series
 - Repeat the process until there is no trend on data ($d \leq 2$)
 - There is a function in R that returns the number of differences needed to achieve stationary

```
ndiffs(x, alpha = 0.05, test = c("kpss", "adf", "pp"), max.d = 2)
```

x is your timeseries

How to Specify Model Orders?


- Step 3: Plot ACF and PACF and look the behavior of non-seasonal lags



If cuts off (MA process)

If tails off, i.e., slow decay (AR process)

ACF gives you the value of q



If cuts off (AR process)

If tails off, i.e., slow decay (MA process)

PACF gives you the value of p

How to Specify Model Orders?

Move to seasonal component


- Step 4: Check if seasonal differencing is needed
 - ▣ Seasonality can usually be verified by plotting your time series or by spikes at equally spaced lags (multiples of lag s) on ACF and PACF
 - ▣ If you are still not sure about the existence of seasonal component you can run statistical tests
 - ▣ The seasonal difference is needed if the seasonal pattern is strong and stable over time
 - ▣ There is also a function in R that returns the number of seasonal differences needed to achieve stationary

```
nsdiffs(x, m = frequency(x), test = c("ocsb", "ch"), max.D = 1)
```


How to Specify Model Orders?


Move to seasonal component

- Step 5: Again ACF and PACF plots but now look at the seasonal lags



Multiple spikes at seas. lag (SAR process)
Single Spike (SMA process)

ACF gives you the value of Q



Multiple spikes at seas. lag (SMA process)
Single Spike (SAR process)

PACF gives you the value of P

Some Rules for ARIMA Modeling

- If the series have positive autocorrelations out to a high number of lags, then it probably needs differencing more than once
- If the lag-1 autocorrelation is zero or negative, or the autocorrelations are all small with no pattern, then the series does not need a higher order of differencing
- If the lag-1 autocorrelation is -0.5 or more negative, the series may be overdifferenced

Some Rules for ARIMA Modeling

- Never use more than 1 order of seasonal differencing
 $D \leq 1$
- Never use more than 2 orders of total differencing
 $\text{seasonal } (D) + \text{nonseasonal } (d) \leq 2$
- If the autocorrelation at the seasonal period is positive, consider adding an **SAR** term to the model
- If the autocorrelation at the seasonal period is negative, consider adding an **SMA**
- Avoid mixing SAR and SMA terms in the same model, and avoid using more than one of either kind

$$P + Q \leq 1$$

P or Q should be always equal to 1.

The ndiffs() and nsdiffs() tests

Augmented Dickey Fuller (ADF)

Check for stochastic trend

$$\begin{cases} H_0: \phi = 1 \text{ (i.e. contain a unit root)} \\ H_1: \phi < 1 \text{ (i.e. is stationary)} \end{cases}$$

Kitawoski-Phillips-Schmidt-Shin (KPSS)

Check for deterministic & stochastic trend

$$\begin{cases} H_0: \text{stationary around determinist trend} \\ H_1: \text{contain a unit root} \end{cases}$$

Phillips-Perron test (PP)

Check for stochastic trend

$$\begin{cases} H_0: \phi = 1 \text{ (i.e. contain a unit root)} \\ H_1: \phi < 1 \text{ (i.e. is stationary)} \end{cases}$$

Canova and Hansen (CH)

Check for seasonal pattern

$$\begin{cases} H_0: \text{seasonal pattern is stable} \\ H_1: \text{not stable.} \end{cases}$$

Osborn,. Chui, Smith and Birchenhall (OCSB)

Check for seasonal unit root

$$\begin{cases} H_0: \text{seasonal unit root.} \\ H_1: \text{no seasonal unit root} \end{cases}$$

Covariates in ARIMA

ARIMAX model

More on ARIMA and SARIMA in R

Constants in the model

Constants in ARIMA models in R

- The function `Arima()` takes three arguments regarding constants
 - ▣ `include.mean` – only has effect when $d = 0$ and is true by default
 - ▣ `include.drift` – allows $\mu \neq 0$ when $d = 1$ and is false by default
 - ▣ `include.constant` – more general. If TRUE, will set `include.mean=TRUE` if $d = 0$ and `include.drift=TRUE` if $d = 1$
 - ▣ Note: For $d > 1$ no constant is allowed [This is differencing for lagging 1.](#)
- The function `auto.arima` automates the inclusion of a constant
 - ▣ If `allowdrift=FALSE` is specified then the constant is only allowed when $d = 0$
- The `arima()` function only has the `include.mean` option



THANK YOU !

luana.marangon.lima@duke.edu

Master of Environmental Management Program
Nicholas School of the Environment - Duke University